

# Ivan Brugere

---

Lead AI Research Scientist - JPMorganChase

ivan@ivanbrugere.com

github.com/ivanbrugere

Chicago, IL

## Current Objective

I am a Lead AI Research Scientist at JP Morgan with 6 years of experience focusing on perturbation-based robustness, adversarial robustness, and privacy-preserving ML. I have 25+ publications, many in top venues, and 12 filed patents. I leverage an interdisciplinary background, domain adaptability, and balanced track record of publications, patents, and deployed systems for high impact research. I am seeking industrial AI scientist roles where I can formulate novel, robust AI methods across large and classical ML models in unique and challenging domains, with an opportunity to mentor and develop emerging researchers.

## Research Statement

My research develops methods to ensure AI systems remain reliable when real-world conditions violate their training assumptions—whether from distributional shifts, adversarial perturbations, or privacy constraints. My work—including my PhD era—has been in diverse, multi-disciplinary domains including biology, ecology, non-profit organizations, finance and legal. I greatly value interdisciplinary partnerships, deeply learning new domain areas, and articulating high-impact computational problems in those areas. I'm seeking collaborative, exploratory research teams and an opportunity to mentor junior scientists and students.

During my current role at JP Morgan, I focused on challenges unique to finance, including explainable and auditable models in highly regulated applications. Furthermore, finance is a high-stake domain, necessitating highly robust methods with respect to model or data perturbation. My PhD (defended 2020) focused on graph topology inference in machine learning; graph learning remains my common lens for formulating computational problems.

## Experience

### **J.P. Morgan Chase & Co. Jan. 2021-Current:**

Lead AI Research Scientist (promoted to **Lead**, Jan. 2025) - Trustworthy AI

Working on AI robustness and fairness: explainable and robust LLM methods, fair agent-based learning, robust and fair tree-based ensembles. Published in ICML, NeurIPS, TMLR, EMNLP, 12 patents filed.

- Co-led research initiative on perturbation robustness tree-based models, leading to 2 NeurIPS publications and 2 patent filings.
- Designed and deployed a model perturbation and auditing system for fairness, targeting firm-wide consumer-facing models (patent filed).
- Developed novel Bayesian model search frameworks for fair mortgage and auto loan pricing (patent filed).
- Led research of two summer AI research interns, yielding 1 NeurIPS, 1 TMLR publication, and 1 successful full-time scientist conversion.

- Co-led external research collaborations with university faculty at NYU, MSU, UMD, Harvard.

### **Salesforce** Jan. 2019–March 2020 Research Scientist - AI For Good

- Collaborated with Salesforce non-profit customers and external collaborators on fairness-aware AutoML applications
- Published research on graph augmentation for equitable access (AAAI AIES 2021), 1 patent filed.

### **Amazon** June–October 2018 Applied Scientist Intern (Mentor: Alex Smola)

Designed deep learning graph APIs on MXNet for scalable training as part of the DGL project. Part of the Amazon Web Services AI Platforms team.

### **Microsoft** May–August 2015 Data Science Intern (Mentor: Marcello Hasegawa)

Developed textual and device models for novelty detection and attribution in the Windows 10 user population.

### **Lawrence Livermore National Laboratory** May–August 2014 Research Intern (Mentor: Brian Gallagher)

Formulated graph inference problems over several scientific research domains.

### **University of Illinois at Chicago** 2013–2015 Research Fellow (Mentor: Prof. Venkat Venkatakrishnan)

Graph-based models for attribute inference and privacy preservation on real mobile device datasets.

### **Technicolor** May–August 2013 Research Intern (Mentor: Fernando Silveira)

Rule discovery for biometric sensor time series data for actionable analysis of film audiences. Developed methods to discover and visualize dynamic audience communities.

### **University of Illinois at Chicago** 2012–2017 Research Assistant (Advisor: Prof. Tanya Berger-Wolf)

Model selection for graph structure inference and prediction. Focusing on ecology and population biology domains.

### **University of Minnesota** 2010–2012 Research Assistant (Advisor: Prof. Vipin Kumar)

Time series change detection and anomaly detection on large remote sensing datasets. Focused on incorporating spatial aspects for change significance testing, and domain-driven information retrieval.

### **University of Minnesota** 2004–2007 Web Applications Developer (Department of Computer Science)

### **University of Minnesota** 2002–2003 Web Applications Developer (College of Liberal Arts)

## Education

### **University of Illinois at Chicago** 2012–2020 Computer Science PhD (Advisor: Prof. Tanya Berger-Wolf) Thesis: Network Structure Inference: Methodology and Applications.

### **University of Minnesota** 2009–2012 Computer Science M.S. (Advisor: Prof. Vipin Kumar) Thesis: Approximate Search on Massive Spatiotemporal Datasets.

### **The New School** 2007–2009 International Affairs M.A.

**University of Minnesota** 2002-2007 Computer Science B.S., Cultural Studies and Comparative Literature B.A.

## Selected Patents

### **System and method for generating constrained loan pricing**

*Automated loan pricing system incorporating regulatory constraints and fairness requirements while optimizing financial objectives.*

**I. Brugere**, M. Hosking, S. Sharma, F. Lecue, Y. Tan, J. Stettler, H. Zhao, P. Glover, D. Kapadia, G. Ciraulo, D. Bollum, D. Magazzeni, L.C. Liang  
US Patent App. 18/397,698, 2025

### **System and method for grounding outputs in tabular generative artificial intelligence**

*Method ensuring LLM-generated table analysis outputs are verifiable and traceable to source data for auditability in regulated applications.*

**I. Brugere**, S. Kariyappa, S. Sharma, F. Lecue, G. Nguyen  
US Patent 12,436,935, 2025

### **System and method for graph-based resource allocation using neural networks**

*Neural network approach for optimizing resource distribution across graph-structured systems, enabling equitable allocation in networked environments.*

G.S. Ramachandran, **I. Brugere**, L. Varshney, C. Xiong  
US Patent 12,165,053, 2024

### **Method and system for improving model fairness by using explainability techniques**

*Framework leveraging model explainability methods to identify and mitigate fairness issues by analyzing feature contributions across demographic groups.*

**I. Brugere**, D. Magazzeni, N. Marchesotti, D. Heike, F. Zhao, E. Wang, H. Shu, M. Gabriel, M. Veloso, C. Tilli, S. Dutta, B. Mallik, A. Onigbanjo  
US Patent App. 17/968,220, 2024

## Selected Publications

### **Calibrating LLM Confidence by Probing Perturbed Representation Stability**

*Novel method applying adversarial perturbations to LLM hidden states to assess confidence via internal stability, reducing calibration error by ~55% across multiple models.*

R. Khanmohammadi, E. Miah, M. Mardikoraem, S. Kaur, **I. Brugere**, C. Smiley, K.S. Thind, M.M. Ghassemi (EMNLP 2025)

### **Cross-Domain Graph Data Scaling: A Showcase with Diffusion Models**

*Introduces UniAug, a universal graph augmentor using discrete diffusion models pre-trained on thousands of graphs to enable cross-domain data scaling and adaptive structure enhancement for downstream tasks.*

W. Tang, H. Mao, D. Dervovic, **I. Brugere**, S. Mishra, Y. Xie, J. Tang (NeurIPS 2025)

### **Interpretable LLM-based Table Question Answering**

*Introduces Plan-of-SQLs (POS) method decomposing table queries into transparent SQL steps for interpretable reasoning, achieving competitive accuracy with 25x fewer LLM calls.*

G. Nguyen, **I. Brugere**, S. Sharma, S. Kariyappa, A.T. Nguyen, F. Lecue (TMLR June 2025)

## RashomonGB: Analyzing the Rashomon Effect and Mitigating Predictive Multiplicity in Gradient Boosting

*First systematic analysis of the Rashomon effect in gradient boosting with novel techniques to explore model sets and reduce predictive multiplicity for fairer model selection.*

H. Hsu, **I. Brugere**, S. Sharma, F. Lecue, C.F. Chen (NeurIPS 2024)

## Comparing Apples to Oranges: Learning Similarity Functions for Data Produced by Different Distributions

*Efficient sampling framework learning cross-group similarity functions with limited expert feedback, enabling fair comparisons of data from different demographic distributions.*

L. Tsepenekas, **I. Brugere**, F. Lecue, D. Magazzeni (NeurIPS 2023)

## Additional Publications (2015-2026)

Full list available at: scholar.google.com/citations?user=JGIGUcsAAAAJ

### Recent Works 2022-2026

#### Revisiting ML Training under Fully Homomorphic Encryption: Convergence Guarantees, Differential Privacy, and Efficient Algorithms

Y. Zhou, M. Liang, **I. Brugere**, D. Dervovic, Y. Guo, A. Polychroniadou, M. Wu, D. Dachman-Soled (in submission)

#### MAFE: Enabling Equitable Algorithm Design in Multi-Agent Multi-Stage Decision-Making Systems

Z.M. Lazri, A. Nakra, **I. Brugere**, D. Dervovic, A. Polychroniadou, F. Huang, D. Dachman-Soled, M. Wu (in submission)

#### How Reliable are Confidence Estimators for Large Reasoning Models? A Systematic Benchmark on High-Stakes Domains

R. Khanmohammadi, E. Miah, S. Kaur, C. Smiley, **I. Brugere**, K.S. Thind, M.M. Ghassemi (EACL 2026)

#### The Unseen Threat: Residual Knowledge in Machine Unlearning under Perturbed Samples

H. Hsu, P. Niroula, Z. He, **I. Brugere**, F. Lecue, C.F. Chen (NeurIPS 2025)

#### Balancing Fairness and Accuracy in Data-Restricted Binary Classification

Z. McBride Lazri, D. Dervovic, A. Polychroniadou, **I. Brugere**, D. Dachman-Soled, M. Wu (ACM TKDD August 2025)

#### Investigating the Temporal Association of Biomedical Research on Small Business Funding: A Bibliometric and Data Analytic Approach

R. Khanmohammadi, S. Kaur, C.H. Smiley, T. Alhanai, **I. Brugere**, A. Nourbakhsh, M.M. Ghassemi (IEEE TCSS)

#### A Canonical Data Transformation for Achieving Inter- and Within-group Fairness

Z. McBride Lazri, **I. Brugere**, X. Tian, D. Dachman-Soled, A. Polychroniadou, D. Dervovic, M. Wu (IEEE TIFS)

#### Bounding the Accuracy Loss for Graphical Model Based Synthetic Data Generation in Privacy-Preserving Machine Learning

Y. Zhou, **I. Brugere**, D. Dachman-Soled, D. Dervovic, M. Liang, A. Polychroniadou, M. Wu, (ICML 2023)

**Hyper-parameter Tuning for Fair Classification without Sensitive Attribute Access**

A.K. Veldanda, **I. Brugere**, S. Dutta, A. Mishler, S. Garg (TMLR)

**Fairness via In-Processing in the Over-parameterized Regime: A Cautionary Tale with MinDiff Loss**

A.K. Veldanda, **I. Brugere**, J. Chen, S. Dutta, A. Mishler, S. Garg (TMLR)

**Earlier Works 2015-2021****Parameterized Explanations for Investor/Company Matching**

S. Kaur, **I. Brugere**, A. Stefanucci, A. Nourbakhsh, S. Shah, M. Veloso (ICAI'21 Workshop on Explainable AI in Finance)

**GAEA: Graph Augmentation for Equitable Access via Reinforcement Learning**

G.S. Ramachandran, **I. Brugere**, L.R. Varshney, C. Xiong (AAAI AIES 2021)

**Evaluation of crowdsourced mortality prediction models as a framework for assessing artificial intelligence in medicine**

T. Bergquist, T. Schaffter, Y. Yan, T. Yu, **I. Brugere** et al. (Journal of the American Medical Informatics Association)

**A continuously benchmarked and crowdsourced challenge for rapid development and evaluation of models to predict COVID-19 diagnosis and hospitalization**

Y. Yan, T. Schaffter, T. Bergquist, T. Yu, J. Prosser, Z. Aydin, A. Jabeer, **I. Brugere**, et al. (JAMA Network Open)

**Network Structure Inference: Methodology and Applications**

**I. Brugere** (Ph.D. Thesis)

**Network Structure Inference, A Survey: Motivations, Methods, and Applications**

**I. Brugere**, B. Gallagher, T. Y. Berger-Wolf (ACM Computing Surveys)

**Network model selection with task-focused minimum description length **I. Brugere**, T.Y. Berger-Wolf**  
(WWW'18: BigNet Workshop on Learning Representations for Big Networks)**Coordination Event Detection and Initiator Identification in Time Series Data**

C. Amornbunchornvej, **I. Brugere**, A. Strandburg-Peshkin, D. Farine, M.C. Crofoot, T.Y. Berger-Wolf (ACM TKDD)

**Evaluating Social Networks Using Task-Focused Network Inference**

**I. Brugere**, C. Kanich, T.Y. Berger-Wolf (KDD'17: Workshop on Mining and Learning in Graphs)

**Both Nearest Neighbours and Long-term Affiliates Predict Individual Locations During Collective Movement in Wild Baboons**

D. Farine, A. Strandburg-Peshkin, T.Y. Berger-Wolf, B. Ziebart, **I. Brugere**, J. Li, M. Crofoot (Nature Scientific Reports)

**Social Information Improves Location Prediction in the Wild**

J. Li, **I. Brugere**, B. Ziebart, T. Y. Berger-Wolf, M. Crofoot, D. Farine (AAAI'15: Workshop on Trajectory-based Behaviour Analytics)

**PhD Scholarships and Awards**

## 2014-2016

- NSF IGERT Electronic Security and Privacy Fellowship
- University of Illinois at Chicago, Chancellor's Graduate Research Fellowship

## Other

- IEEE ICDM 2017 Travel Award
- SIAM SDM 2017 Travel Award
- 2016 ACM Tapia Celebration of Diversity in Computing, Travel Award
- 2016 ACM SIGKDD Broadening Participation in Data Mining Travel Award
- 2016 ACM WSDM Travel Award
- 2015 IEEE ICDM Travel Award
- 2015 ACM Ubicomp Broadening Participation Travel Award
- 2015 ACM SIGKDD Ram Kumar Memorial Travel Award
- 2015 SIAM CSE Travel Award supported by the Sustainable Horizons Institute
- Fifty for the Future Award supported by the Illinois Technology Foundation
- 2014 Google Lime Scholarship
- 2014 ACM BCB Travel Award
- 2014 ACM SIGKDD Broadening Participation in Data Mining Travel Award
- 2014 ACM Tapia Celebration of Diversity in Computing, Travel Award

## Community Activities

### Workshop organization (2016-2023)

- NLP and Network Analysis in Financial Applications (ACM ICAIF'23)
- PhD Forum (IEEE ICDM'19)
- NetInf'17: First Workshop on Inferring Networks from Non-Network Data (SIAM SDM'17)
- Inferring Networks from Non-Network Data (SIAM AM'16)

### PC Member/Reviewer (2018-Present)

- Conferences: AAAI, CIKM, FAccT, ICDM, ICLR, IJCAI, KDD, NeurIPS, PAKDD, SDM, TheWebConf, WSDM
- Journals: ACM CSUR, IEEE TKDE, ACM TKDD, KAIS

### Tutorials (2018)

Modeling Data with Networks + Network Embedding: Problems, Methodologies and Frontiers

**I. Brugere**, B. Perozzi, P. Cui, W. Zhu, J. Pei, T.Y. Berger-Wolf (KDD 2018)

### Other service (2014-2020)

- ACM Tapia Celebration of Diversity in Computing 2020 Plenary Speaker
- ACM Tapia Celebration of Diversity in Computing 2020 Accessibility Committee
- Bloomberg Data For Good Exchange Program Committee
- Google Lime campus ambassador
- University of Washington-AccessSTEM volunteer

- ACM SIGKDD Broadening Participation in Data Mining Coordinator and Mentoring Co-Chair (2014, 2016, 2017)

## Teaching (2017)

Teaching Assistant: Computer Algorithms I (Senior-level), University of Illinois at Chicago.

## Technical Skills

**Languages:** Python, Julia, Scala **AI/ML:** PyTorch, scikit-learn, XGBoost, Gymnasium, LM Studio **Graph Learning:** DGL, PyTorch Geometric, NetworkX **Tools:** Git, Docker, AWS, Jupyter, VS Code

## Links

- [ivan@ivanbrugere.com](mailto:ivan@ivanbrugere.com)
- [github.com/ivanbrugere](https://github.com/ivanbrugere)
- [linkedin.com/in/ivanbrugere](https://linkedin.com/in/ivanbrugere)
- [scholar.google.com/citations?user=JGlGUcsAAAAJ](https://scholar.google.com/citations?user=JGlGUcsAAAAJ)
- [orcid.org/0000-0002-2953-3746](https://orcid.org/0000-0002-2953-3746)